

# CSC14120 – PARALLEL PROGRAMMING

## FINAL PROJECT

### 1. Introduction

#### 1.1 Problem Statement

Feature engineering is a fundamental challenge in machine learning: how do we automatically discover good representations of data that capture its underlying structure? In this project, you will implement an **Autoencoder-based unsupervised feature learning system** for image classification on the **CIFAR-10 dataset**.

Traditional supervised learning trains a model end-to-end using labelled examples. However, autoencoders take a different approach: they learn meaningful representations by attempting to reconstruct the input data itself, without requiring any labels during the feature learning phase. This unsupervised pre-training can discover features that are often more robust and generalizable than those learned through direct supervision alone.

**Your task** is to build and optimize a complete two-stage pipeline:

#### Stage 1 - Unsupervised Feature Learning:

- Train a convolutional autoencoder to reconstruct CIFAR-10 images
- The autoencoder learns to compress  $32 \times 32 \times 3$  images into a compact 1,024-dimensional representation
- No labels are used during this training phase
- The network learns features that capture important visual patterns (edges, textures, shapes)

#### Stage 2 - Supervised Classification:

- Extract features from the trained encoder for all images
- Train an SVM classifier on these learned features with class labels
- Evaluate classification performance on the test set

The primary focus is on **implementing and optimizing the autoencoder training and inference in CUDA**, while using existing libraries (LIBSVM) for the SVM classifier. Through systematic optimization, you will accelerate the autoencoder from taking hours on CPU to completing in seconds on GPU.

#### 1.2 CIFAR-10 Dataset

**CIFAR-10** is one of the most widely used benchmark datasets in computer vision and machine learning research. It provides a challenging yet computationally manageable image classification task.

### Dataset Specifications:

- **Image size:** 32×32 pixels (RGB)
- **10 classes:** airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck
- **Training set:** 50,000 images (5,000 per class)
- **Test set:** 10,000 images (1,000 per class)
- **Total images:** 60,000
- **Format:** Binary files with uint8 pixel values



Figure 1 - Examples from CIFAR-10 Dataset

### Dataset Organization for This Project:

- **Autoencoder training:** All 50,000 training images (labels ignored)
- **SVM training:** Same 50,000 images with labels (using extracted features)
- **Evaluation:** 10,000 test images with labels

**Dataset Link:** <https://www.cs.toronto.edu/~kriz/cifar.html>

### 1.3 Expected Performance Targets:

Metric	Target
Autoencoder training time	<10 minutes
Feature extraction time	<20 seconds for all 60K images
Test classification accuracy	60-65%

Metric	Target
GPU speedup over CPU	>20x

## 2. Background knowledge

### 2.1 Autoencoders for Unsupervised Feature Learning

#### What is an Autoencoder?

An autoencoder is a neural network trained to reconstruct its input, forcing the network to learn a compressed, meaningful representation in the process. It consists of two parts:

- **Encoder:** Compresses the input into a low-dimensional latent representation (feature vector)
- **Decoder:** Reconstructs the original input from the latent representation

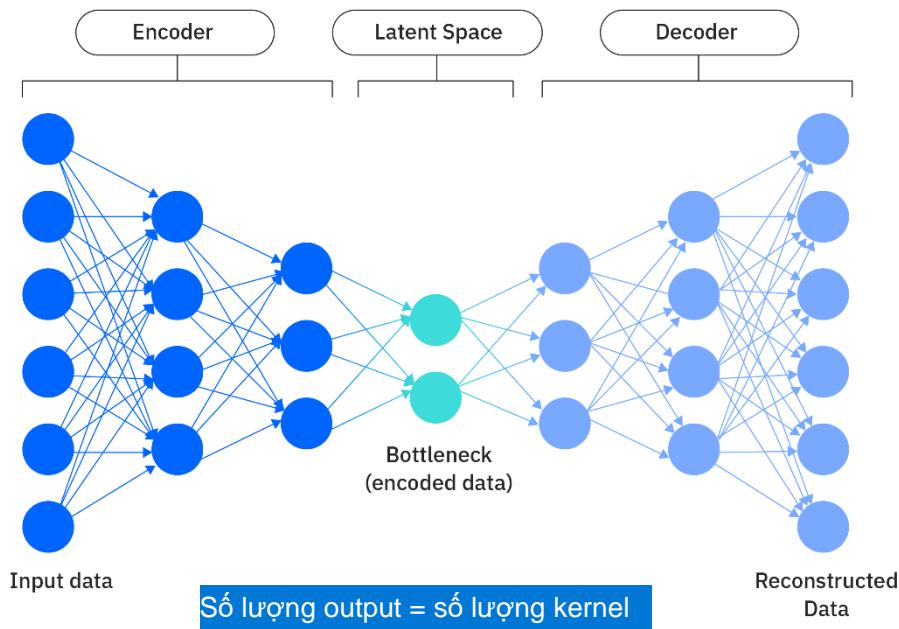


Figure 2 - The architecture of a standard autoencoder neural network.

#### Training Objective:

$$\begin{aligned} \text{Loss} &= \text{MSE}(\text{Input}, \text{Reconstructed\_Output}) \\ &= (1/N) * \sum (x - \text{decoder}(\text{encoder}(x)))^2 \end{aligned}$$

The network learns to minimize reconstruction error, forcing the encoder to capture essential information about the input

#### Key Concepts:

### 1. Bottleneck Layer (Latent Space):

- The smallest layer in the middle
- Forces compression and feature extraction
- In this project:  $(6, 6, 128) = 4,608$  dimensions

### 2. Symmetric Architecture:

- Decoder mirrors the encoder
- Up sampling operations reverse the down sampling
- Helps in reconstruction quality

### 3. Feature Extraction:

- After training, we discard the decoder
- Use only the encoder to extract features
- These features feed into the SVM classifier

### Recommended Reading Materials:

- **Foundational Papers:**

Hinton, G. E., & Salakhutdinov, R. R. (2006). "Reducing the Dimensionality of Data with Neural Networks." *Science*, 313(5786), 504-507.

Link: <https://www.cs.toronto.edu/~hinton/science.pdf>

- **Book Chapters:**

Goodfellow, I., Bengio, Y., & Courville, A. "Deep Learning" (2016)

Chapter 14: Autoencoders

Link: <https://www.deeplearningbook.org/contents/autoencoders.html>

- **Video guide:** Autoencoders and Representation Learning

Link: <https://www.youtube.com/watch?v=R3DNKE3zKFk>

- **Reference source code:**

Link: <https://github.com/turkdogan/autoencoder>

Link: <https://github.com/tbennun/cudnn-training>

- **CNN references:** Since your autoencoder uses convolutional layers, knowledge of CNNs is required

- ✓ This video "[How Convolutional Neural Networks work](#)" give a very good explanation with respect to CNNs. You can also check other good materials about Neural Network [here](#) from the same authors.
- ✓ This Standford [cheatsheet](#) give an excellent explain on how CNNs works
- ✓ This "[Dive into deep learning](#)" book also give through explain on CNN.

## 2.2 Support Vector Machine (SVM) for Classification

**Your Role with SVM:** You do NOT need to implement SVM from scratch. Use existing optimized libraries.

### Recommended SVM Libraries:

#### 1. LIBSVM (Primary Recommendation)

- Most popular SVM library
- Easy to use C/C++ interface
- Supports multi-class classification
- RBF kernel built-in
- Link: <https://www.csie.ntu.edu.tw/~cjlin/libsvm/>
- Documentation: <https://github.com/cjlin1/libsvm>

#### 2. LIBLINEAR (Alternative for Linear SVM)

- Faster for linear kernels
- Good if you want to experiment
- Link: <https://www.csie.ntu.edu.tw/~cjlin/liblinear/>

#### 3. ThunderSVM (GPU-Accelerated)

- GPU implementation of SVM
- Optional: Could explore for extra credit
- Link: <https://github.com/Xtra-Computing/thundersvm>

## 3. Project Pipeline

The project follows a clear 5-step pipeline:

### Step 1: Load CIFAR-10 Images

- ✓ 50,000 training images (for autoencoder)
- ✓ 10,000 test images
- ✓ Preprocess: normalize to [0,1]

### **Step 2: Train Autoencoder (Your CUDA Code)**

- ✓ Use all 50k images (unsupervised training)
- ✓ Ignore labels during autoencoder training
- ✓ CNN-based autoencoder architecture (see Section 4)
- ✓ Train to minimize reconstruction loss
- ✓ Save encoder weights

### **Step 3: Extract Features (Your CUDA Code)**

- ✓ Load trained encoder weights
- ✓ Run encoder forward pass (no decoder)
- ✓ train\_features: (50000, 1024)
- ✓ test\_features: (10000, 1024)

### **Step 4: Train SVM (Library)**

- ✓ Input: train\_features + labels
- ✓ Kernel: RBF (Radial Basis Function)
- ✓ Hyperparameters: C=10, gamma=auto
- ✓ Output: trained SVM model
- ✓ (*Can use other SVM setup as you see fit*)

### **Step 5: Evaluate**

- ✓ Predict on test\_features using SVM
- ✓ Calculate accuracy, confusion matrix
- ✓ Expected accuracy: 60-65%
- ✓ Compare with baseline methods

### **Key Points:**

- ✓ Your focus: Steps 2 and 3 (autoencoder implementation in CUDA)
- ✓ Using libraries: Step 4 (SVM training)
- ✓ Evaluation: Step 5 (measure success)

## **4. Network Architecture**

### **4.1 Architecture Overview**

INPUT: (32, 32, 3) → ENCODER (compress) → LATENT: (8, 8, 128) = 1,024 features → DECODER (reconstruct) → OUTPUT: (32, 32, 3)

### **4.2 Detailed Architecture Specification**

#### **ENCODER (Down sampling Path)**

INPUT: (32, 32, 3)
↓

```

Conv2D(256 filters, 3x3 kernel, padding=1, stride=1) + ReLU
→ (32, 32, 256)
↓
MaxPool2D(2x2, stride=2)
→ (16, 16, 256)
↓
Conv2D(128 filters, 3x3 kernel, padding=1, stride=1) + ReLU
→ (16, 16, 128)
↓
MaxPool2D(2x2, stride=2)
→ (8, 8, 128)
↓
LATENT REPRESENTATION: (8, 8, 128) = 1,024 dimensions

```

## DECODER (Upsampling Path - Mirror of Encoder):

```

LATENT: (8, 8, 128)
↓
Conv2D(128 filters, 3x3 kernel, padding=1, stride=1) + ReLU
→ (8, 8, 128)
↓
UpSample2D(2x2) [Nearest neighbor or bilinear]
→ (16, 16, 128)
↓
Conv2D(256 filters, 3x3 kernel, padding=1, stride=1) + ReLU
→ (16, 16, 256)
↓
UpSample2D(2x2) [Nearest neighbor or bilinear]
→ (32, 32, 256)
↓
Conv2D(3 filters, 3x3 kernel, padding=1, stride=1) [No activation]
→ (32, 32, 3)
↓
OUTPUT: (32, 32, 3)

```

Or keras setup for reference:

```

input_size = (32, 32, 3)

input_image = Input(shape=input_size)

# Encoder
x = Conv2D(256, (3, 3), activation='relu', padding='same')(input_image)
x = MaxPooling2D((2, 2), padding='same')(x)
x = Conv2D(128, (3, 3), activation='relu', padding='same')(x)
encoded = MaxPooling2D((2, 2), padding='same', name='encoded_layer')(x)

# Decoder
x = Conv2D(128, (3, 3), activation='relu', padding='same')(encoded)
x = UpSampling2D((2, 2))(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same')(x)
x = UpSampling2D((2, 2))(x)

```

```
decoded = Conv2D(3, (3, 3), padding='same')(x)
```

Parameter summary:

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	(None, 32, 32, 3)	0
conv2d_1 (Conv2D)	(None, 32, 32, 256)	7168
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 256)	0
conv2d_2 (Conv2D)	(None, 16, 16, 128)	295040
encoded_layer (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_3 (Conv2D)	(None, 8, 8, 128)	147584
up_sampling2d_1 (UpSampling2D)	(None, 16, 16, 128)	0
conv2d_4 (Conv2D)	(None, 16, 16, 256)	295168
up_sampling2d_2 (UpSampling2D)	(None, 32, 32, 256)	0
conv2d_5 (Conv2D)	(None, 32, 32, 3)	6915
<hr/>		
Total params:	751,875	
Trainable params:	751,875	
Non-trainable params:	0	

## 5. Implementation Guideline

You will implement this project in at least **4 progressive phases**, focusing on CUDA optimization for the autoencoder training and inference.

### Phase 1: CPU Baseline & Data Pipeline

**Objective:** Set up the project infrastructure and create a working CPU baseline.

**What to Implement:**

#### 1.1 Data Loading and Preprocessing

- Create a CIFAR10 Dataset class to handle data loading
- Read CIFAR-10 binary files (5 training batches + 1 test batch)
- Parse the binary format: 1 byte label + 3,072 bytes image per record
- Convert uint8 pixel values [0, 255] to float [0, 1] for normalization
- Implement batch generation for training

- Add data shuffling capability
- Organize train images (50,000), test images (10,000), and their labels in memory

## 1.2 CPU Neural Network Layers

Implement CPU versions of all necessary operations:

- **Convolution (Conv2D):** Apply  $3 \times 3$  kernels with padding and stride
- **ReLU Activation:** Element-wise  $\max(0, x)$  operation
- **Max Pooling:**  $2 \times 2$  pooling to downsample by half
- **Upsampling:** Nearest neighbor interpolation to double spatial dimensions
- **MSE Loss:** Mean squared error between output and target for reconstruction

## 1.3 Autoencoder

- Create an Autoencoder class to encapsulate the network
- Allocate memory for all weight matrices and biases (5 conv layers total)
- Implement weight initialization
- Allocate memory for intermediate activations between layers
- Implement forward pass: sequentially apply encoder layers, then decoder layers
- Implement backward pass:
- Implement feature extraction method: run encoder only and return latent representation
- Add weight saving/loading functionality for model persistence

## 1.4 Training

- Set hyperparameters: batch size (32), epochs (20), learning rate (0.001)
- Loop over epochs and batches
- For each batch: forward pass → compute loss → backward pass → update weights
- Track and display training loss
- Measure time per epoch
- Save trained model weights after completion

### Deliverables:

- Working data loading and preprocessing pipeline

- Complete CPU implementation of all layers
- Baseline performance measurements

## Phase 2: Naive GPU Implementation

**Objective:** Port all operations to GPU with basic parallelization.

### What to Implement:

#### 2.1 GPU Memory Management

- Create GPUAutoencoder class for GPU operations
- Allocate device memory for all weight matrices
- Allocate device memory for all activation buffers
- Allocate device memory for gradients
- Implement functions to copy weights between host and device
- Implement proper memory cleanup (cudaFree)

#### 2.2 Naive GPU Kernels

Design basic GPU kernels where each thread handles simple work:

##### Convolution Kernel:

- Each thread computes one output pixel
- Thread performs nested loops over kernel and input channels
- Uses global memory for all reads/writes
- Handles boundary conditions with padding

##### ReLU Kernel:

- Each thread processes one element
- Simple in-place operation:  $x = \max(0, x)$

##### MaxPooling Kernel:

- Each thread computes one output element
- Finds maximum value in  $2 \times 2$  input window
- Writes result to output

##### Upsampling Kernel:

- Each thread computes one output pixel

- Maps output coordinates back to input coordinates (divide by 2)
- Copies input value to output (nearest neighbour)

### MSE Loss Kernel:

- Parallel reduction to sum squared differences
- Use shared memory for partial sums
- Use atomicAdd for final accumulation

## 2.3 GPU Forward Pass

- Copy input batch from host to device
- Launch kernels sequentially for each layer
- Synchronize after kernel launches
- Copy final output back to host
- Compute loss value

## 2.4 GPU Backward Pass

Design gradient computation kernels:

- Implement gradient of each layer type (conv, relu, maxpool, upsample)
- Backpropagate errors from output to input
- Compute gradients with respect to weights
- Implement simple SGD weight update:  $\text{weight} -= \text{learning\_rate} \times \text{gradient}$

## 2.5 GPU Training Loop

- Set larger batch size for GPU (64 instead of 32)
- For each batch: copy to device → forward → compute loss → backward → update
- Use GPU timer for accurate timing
- Display progress and loss values
- Save trained weights

### Deliverables:

- All layers successfully ported to GPU
- Working GPU training loop

- Forward and backward passes implemented
- Correctness verification against CPU results

### **Phase 3: Advanced Optimization**

Objective: Optimize memory access patterns and apply advanced techniques for maximum performance (May have several versions here)

**Optimization Ideas** (You don't need to implement all, this list is just a suggestion)

#### **Category 1: Memory Optimization**

##### **1. Shared Memory Tiling for Convolution**

Load input tiles cooperatively into shared memory to reduce global memory accesses. Each thread block processes a tile, reusing loaded data across multiple threads.

##### **2. Convert to Matrix Multiplication**

Transform convolution into matrix multiplication, then leverage optimized via matrix multiplication.

##### **3. Memory Coalescing Optimization**

Ensure threads in a warp access consecutive memory address for maximum bandwidth. Reorganize data access patterns and consider changing memory layout

##### **4. Constant Memory for Small Weights**

Store small, frequently accessed data (biases, small kernel weights) in constant memory for fast broadcast access. All threads in a warp can read the same value in a single instruction.

##### **5. Pinned (Page-Locked) Memory**

Use cudaMallocHost for pinned memory on CPU side to enable faster asynchronous transfers. Reduces H2D and D2H transfer time compared to pageable memory.

##### **6. Unified Memory Management**

Use CUDA Unified Memory to simplify memory management and allow on-demand data migration. System automatically handles data movement between CPU and GPU.

##### **7. Memory Pool/Reuse Strategy**

Reallocate all necessary buffers once and reuse them across batches/epochs. Eliminates repeated cudaMalloc/cudaFree overhead which can be significant.

#### **Category 2: Kernel-Level Optimization**

##### **8. Kernel Fusion (Conv + ReLU + Bias)**

Combine multiple operations into single kernel to eliminate intermediate global memory writes/reads. Reduces memory bandwidth requirements and improves arithmetic intensity.

## **9. Block-Level Fusion (Entire Encoder/Decoder)**

Fuse entire encoder or decoder paths into single mega-kernel using shared memory for intermediates. Dramatically reduces global memory traffic but requires careful shared memory management.

## **10. Loop Unrolling**

Unroll loops in convolution, pooling, and other kernels to reduce loop overhead. Compiler can often optimize better with unrolled loops, especially for small kernel sizes.

## **11. Vectorized Memory Access (float4)**

Load/store 4 floats at once using float4 for better memory bandwidth utilization. Particularly effective when processing channels or batch dimensions.

## **12. Optimized Thread Block Dimensions**

Tune block sizes (e.g.,  $16 \times 16$ ,  $32 \times 8$ ) based on profiling to maximize occupancy and minimize wasted threads. Different layers may benefit from different configurations.

## **13. Mixed Precision Training (FP16/FP32)**

Use FP16 for forward pass and most computations, FP32 for weight updates and accumulators. Reduces memory bandwidth.

## **Category 3: Parallelism & Concurrency**

### **14. Gradient Checkpointing**

Store only subset of activations during forward pass, recompute others during backward pass. Trades computation for memory, enabling larger batch sizes.

### **15. Multi-Stream Pipeline**

Use multiple CUDA streams to overlap H2D transfer, computation, and D2H transfer. Can hide transfer latency by processing next batch while current batch trains.

### **16. Batched Operations**

Process multiple images in parallel rather than one at a time. Amortizes kernel launch overhead and improves GPU occupancy.

## **Phase 4: SVM Integration and analysed result**

**Objective:** Complete the pipeline and thoroughly evaluate performance.

## 6. Project Report

### Report Format and Structure

#### Submission Format

Your project report must be submitted as a **Notebook (.ipynb)** that can run on **Google Colab**.

#### Section 1: Problem Description

##### What to Include:

###### 1. Problem Statement

- Clearly define the image classification task
- Describe the motivation for GPU acceleration

###### 2. CIFAR-10 Dataset Overview

- Dataset specifications (size, classes, split)
- Show sample images from each class (use visualization)
- Explain data preprocessing steps (normalization, format)

###### 3. Autoencoder Architecture

- Describe the network architecture with a diagram
- Specify layer dimensions and transformations
- Explain the encoder-decoder structure and latent representation
- Include architecture visualization

###### 4. Project Objectives

- Performance goals (training time, speedup targets, accuracy)
- Technical learning objectives
- Success criteria

#### Section 2: Implementation Phases

For each phase, provide a structured description following this template:

##### Phase 2.1: CPU Baseline Implementation

###### Objectives:



- What you aimed to achieve in this phase
- Why this phase is necessary

### **Implementation Details:**

- **Data Pipeline:** How you loaded and pre-processed CIFAR-10 data
- **Layer Implementations:** Brief description of each layer (Conv2D, ReLU, MaxPool, Upsample)
- **Training Loop:** How you structured the training process
- **Key Code Snippets:** Show 2-3 critical functions (e.g., convolution function signature and main loop structure)

### **Results:**

- Training time per epoch and total training time
- Final reconstruction loss
- Sample reconstructed images (show original vs reconstructed)
- Memory usage

### **Key Takeaways:**

- What did you learn about the algorithm?
- What insights guided your GPU implementation?

---

## **Phase 2.2: GPU Basic Implementation**

### **Objectives:**

- Port CPU code to GPU with basic parallelization
- Verify correctness of GPU kernels
- Establish baseline GPU performance

### **Implementation Details:**

- **Parallelization Strategy:** How you mapped operations to GPU threads
- **Kernel Designs:**
  - Convolution kernel: thread-to-output mapping
  - Pooling kernel: how threads handle  $2 \times 2$  windows
  - Other kernels (ReLU, upsampling)

- **Memory Management:** Device memory allocation strategy
- **Key Code Snippets:** Show kernel signatures and launch configurations

#### **Results:**

- Training time per epoch and total training time
- Speedup over CPU baseline (include table and chart)
- GPU memory usage
- Verification that outputs match CPU (show error metrics)

#### **Profiling Analysis:**

- Basic profiling results (time spent in each kernel type)
- Memory bandwidth utilization (if measured)
- Initial bottleneck identification

#### **Key Takeaways:**

- What was surprisingly fast or slow?
  - Where do you see optimization opportunities?
- 

### **Phase 2.3: GPU Optimized Implementation - Version 1**

**Optimization Focus:** (e.g., Memory Optimization)

#### **Objectives:**

- What specific optimization(s) you targeted
- Expected performance improvement

#### **Implementation Details:**

- **Optimization Technique(s) Applied:**
  - Detailed explanation of the optimization (e.g., shared memory tiling)
  - Why this optimization should help
  - Implementation approach
- **Key Code Snippets:** Show the optimized kernel or key changes

#### **Results:**

- Training time comparison with previous version

- Speedup over previous phase (incremental and cumulative)
- Performance metrics (bandwidth utilization, occupancy)
- Profiling comparison: before vs after

#### **Analysis:**

- Why did this optimization work (or not work as expected)?
- What did profiling reveal?
- What's the next bottleneck?

#### **Key Takeaways:**

- Lessons learned from this optimization
  - Applicability to other problems
- 

### **Phase 2.4: GPU Optimized Implementation - Version 2 (if applicable)**

**Optimization Focus:** (e.g., Kernel Fusion and Advanced Techniques)

Follow the same structure as Version 1:

- Objectives
- Implementation Details
- Results
- Analysis
- Key Takeaways

**Note:** You may have multiple optimization versions. Create a separate subsection for each major optimization iteration.

---

### **Phase 2.5: SVM Integration**

#### **Objectives:**

- Extract features using trained encoder
- Train SVM classifier on learned features
- Evaluate end-to-end classification performance

#### **Implementation Details:**

- **Feature Extraction:** How you extracted 1,024-dim features from encoder
- **LIBSVM Integration:** How you interfaced with LIBSVM
- **Hyperparameter Selection:** SVM parameters chosen (C, gamma, kernel)
- **Key Code Snippets:** Feature extraction and SVM training code

### **Results:**

- Feature extraction time (50K train + 10K test)
- SVM training time
- Classification accuracy on test set
- Per-class accuracy breakdown (table)
- Confusion matrix (visualization)
- Comparison with baseline methods (if available)

### **Analysis:**

- Which classes are easiest/hardest to classify?
- What does the confusion matrix reveal?
- How does accuracy compare to expectations?

### **Key Takeaways:**

- Quality of learned features
- Effectiveness of two-stage approach

---

## **Section 3: Comprehensive Performance Analysis**

### **3.1 Performance Comparison Across All Phases**

Create comprehensive comparison including:

#### **Table Format:**

Phase	Training Time	Speedup (vs CPU)	Incremental Speedup	Memory Usage	Key Optimization
CPU Baseline	1800s	1.0×	-	-	-

Phase	Training Time	Speedup (vs CPU)	Incremental Speedup	Memory Usage	Key Optimization
GPU Basic	180s	10.0×	10.0×	2.1 GB	Parallelization
GPU Opt v1	45s	40.0×	4.0×	2.3 GB	Shared memory
GPU Opt v2	25s	72.0×	1.8×	2.5 GB	Kernel Fusion + Streams

### Visualization Requirements:

- Bar chart showing training time across phases
- Line graph showing cumulative speedup

## Section 4: Lessons Learned and Challenges Overcome

### 4.1 Key Technical Insights

What you have learned about: CUDA Programming, Deep Learning, Performance Optimization

### 4.2 Major Challenges and Solutions

Present 2-3 significant challenges using this format:

- ✓ Challenge 1: [Brief Title]
- ✓ Problem: One sentence describing the issue.
- ✓ Solution: 1-2 sentences explaining how you solved it.
- ✓ Lesson: One sentence on what you learned.

## Section 5: Conclusion and Future Work

### 5.1 Project Summary

- Recap of what was accomplished
- Final performance metrics summary table
- Achievement of original objectives

### 5.2 Key Achievements

Highlight your best results:

- Maximum speedup achieved
- Classification accuracy

- Most successful optimization
- Technical skills mastered

### 5.3 Limitations

Honestly discuss:

- Current performance bottlenecks
- Accuracy limitations
- Implementation constraints

### 5.4 Future Improvements

## 7. Project Deliverable

### Submission Requirements

You must submit the following items:

#### 1. Team Plan and Work Distribution

- Document detailing each team member's responsibilities
- Task breakdown and timeline
- Contribution percentage for each member

#### 2. Project Report

- **Jupyter Notebook (.ipynb)** as specified in Report Requirements
- Must be executable on Google Colab
- Export to PDF for archival purposes

#### 3. Source Code Package

- **All source code files** (.cpp, .cu, .h, header files, etc.)
- **README.md** with:
  - Setup instructions (dependencies, CUDA version, libraries)
  - Compilation commands
  - Execution instructions with example commands
  - Hardware requirements (GPU model, memory)
  - Expected outputs



- **Trained model weights** (if file size permits, otherwise provide download link)

#### **4. Presentation Video (15-20 minutes)**

- **Upload to YouTube with Unlisted visibility**
- **Include link** in your notebook report and README file
- **DO NOT upload video to Moodle** (file size restrictions)

#### **Video Content Should Cover:**

- Problem overview and approach (2-3 min)
- Live demonstration of running code (5-7 min)
- Performance results and analysis (5-7 min)
- Key optimizations and lessons learned (3-5 min)